

1. Did your estimations reflect the actual effort?

Yeah, in both the sprints, the 8-point stories (Generate Weekly Meal Plan and Modify Delivery Schedule) produced the largest single drop.

Small stories such as Stop Subscription (2 points) and View Subscriber Profile (2 points) were completed quickly.

The estimates are relative rather than exact, and safety-critical work like the allergen check probably carries more uncertainty than 5 points.

2. Was your backlog well-prioritized?

Yes. Foundational and safety-critical stories were placed in Sprint 1 because everything else depends on them. The dietitian review workflow and subscription and delivery management followed in Sprint 2, once a meal plan could exist, and lower-impact stories like Stop Subscription were rated Medium and done last.

Splitting the 8-point stories into smaller ones would allow more flexible re-prioritization.

3. How did your simulated sprint align with your plan?

Both sprints delivered everything committed, with 6 stories of 28 points in Sprint 1 and 8 stories of 32 points in Sprint 2, completed in dependency order.

4. What insights did the burndown chart give about your team's capacity?

Both sprints burned down to zero before their end times, so the full commitment was completed with time to spare, in 28 points (Sprint 1) and 32 points (Sprint 2).

This suggests a capacity of roughly 30 points per sprint, with room to take on slightly more.

The step sizes also show that large stories (5 and 8 points) cause big sudden drops, while smaller stories give a smoother chart, so breaking work into smaller pieces would make progress easier to track.

Sprint 2's flat guideline and jump at the start show how scope added after sprint start distorts the chart.